

Challenges faced and lessons learnt while building a multi-cloud Data platform

Antonio Murgia – Big Data conference EU 2023

Who am I

 Data Architect at AgileLab.

I spent the last 7 years first implementing, then leading teams and ultimately designing from the ground-up scalable streaming and batch data solutions.

I'm an open source contributor 
(Apache Spark ⚡ and Project Nessie 🦕)

In my free time I love to hike  and climb  mountains  and I'm up for playing any kind of sport 



Great place to work 😎



witboost

Great
Place
To
Work.[®]

Born as a Data management consultancy company in 2013, based in Turin (Italy)
Focused on scalable technologies and data practices – which is a way to scale too

We are now 150 engineers all over europe!

In recent years we started to develop witboost, a product that helps big organizations build their internal data development platform based on their own needs, processes, practices and policies

⚠️ DISCLAIMER ⚠️

This is not a guide on «how to» implement a multi-cloud data platform.

We were NOT implementing a «greenfield» data platform.

**This is the story of how we managed to transform
an AWS centric platform to a multi-cloud one.**

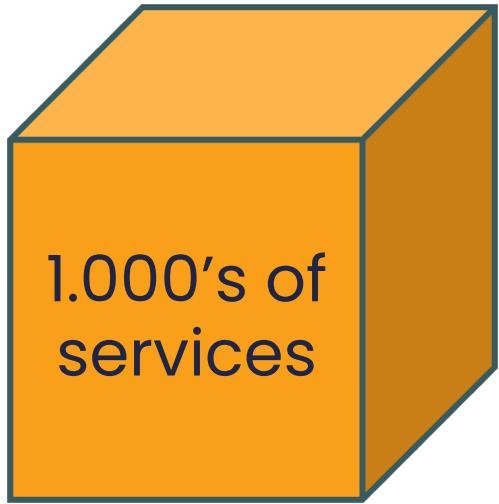
Once upon a time

We started building a service + data platform on AWS in 2019
The platform follows an “holistic” approach:

domain teams  build,  deploy and  operate
in the same platform both operational and data services/products

In retrospective we were building a very prehistoric kind of Data Mesh +1 year before Zhamak blogged about it on Martin Fowler website

Context



Technology Stack



Amazon EKS

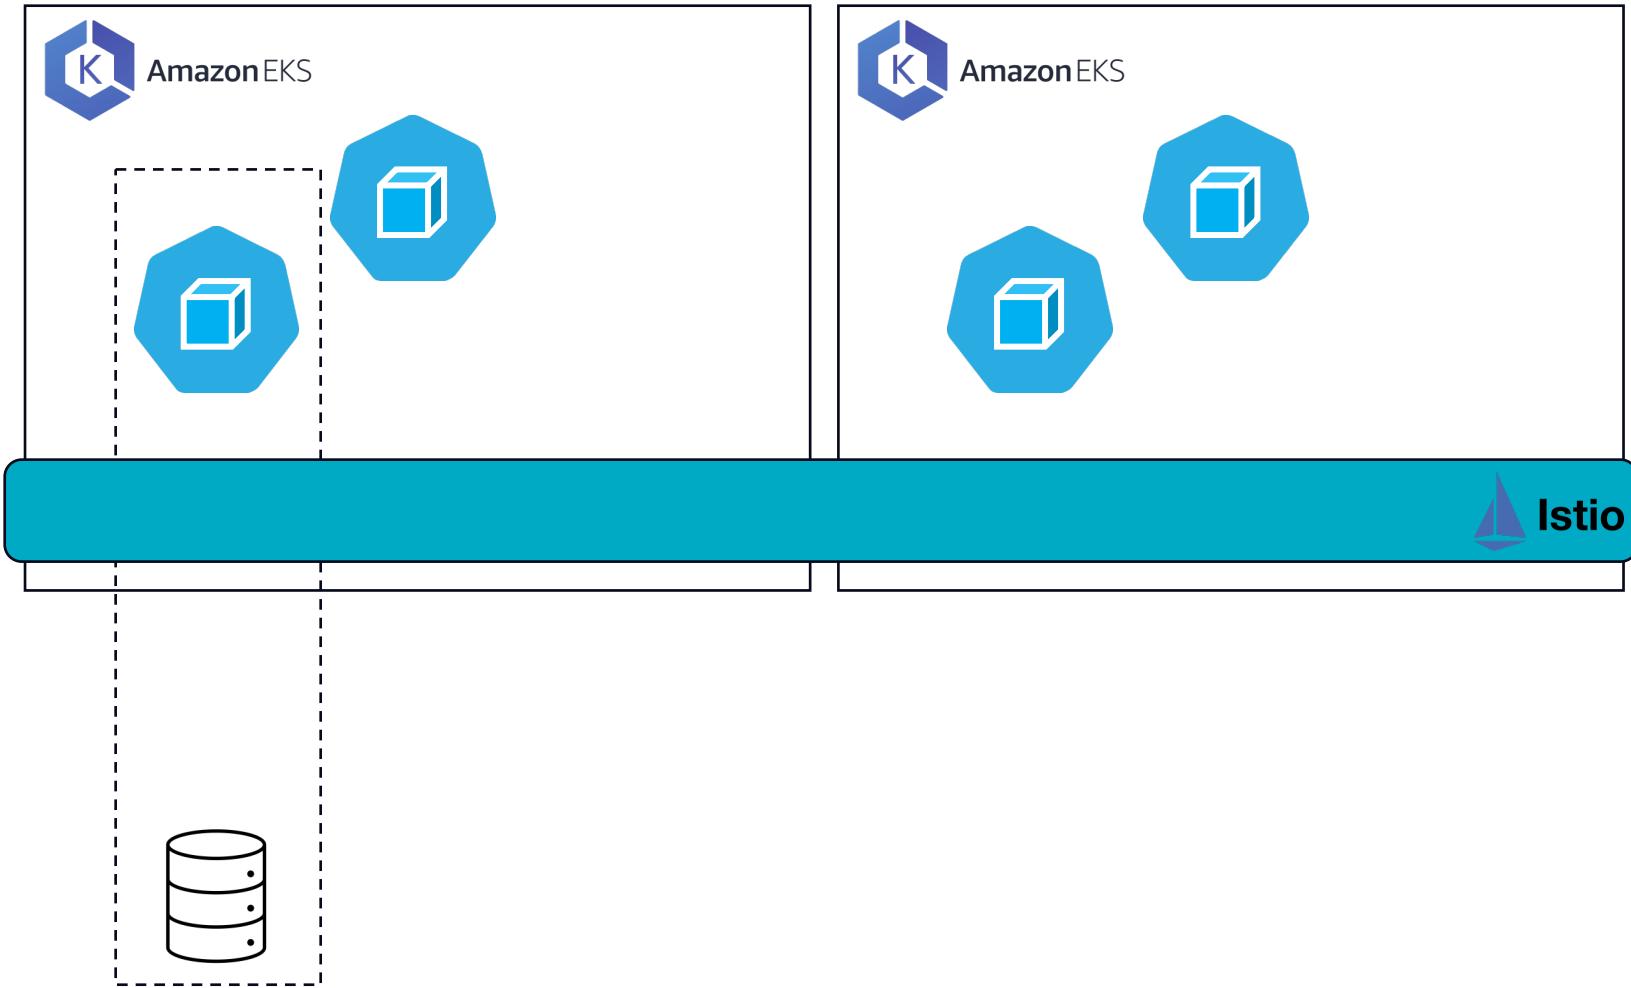


Amazon EKS

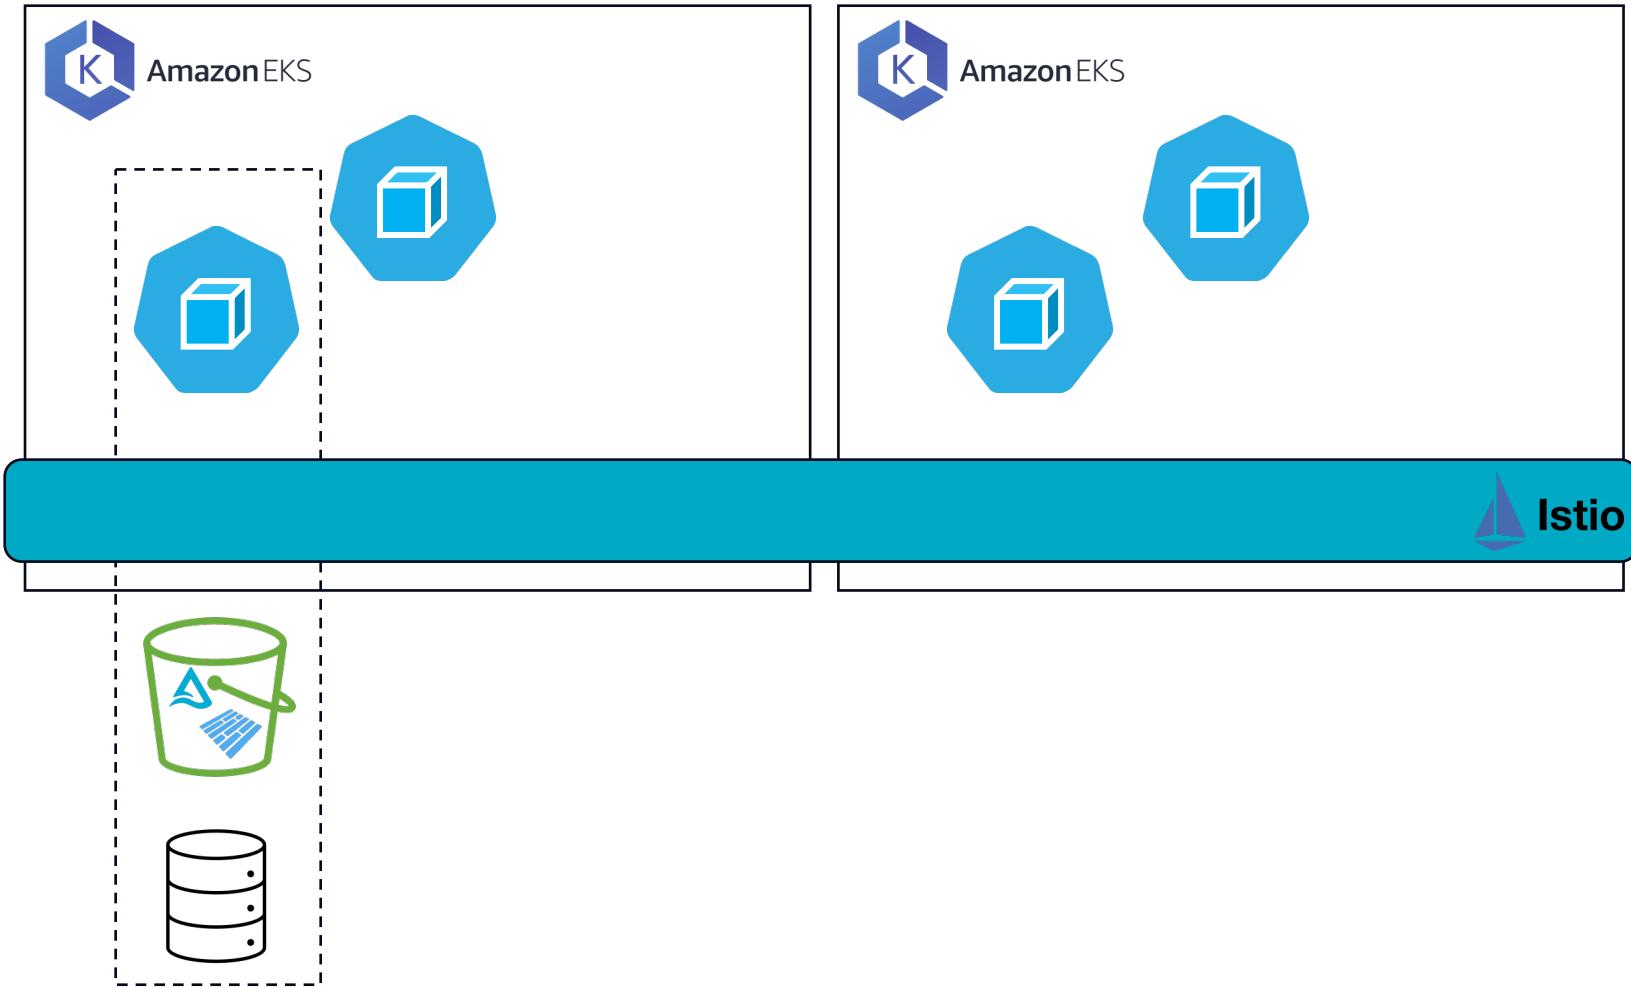


Istio

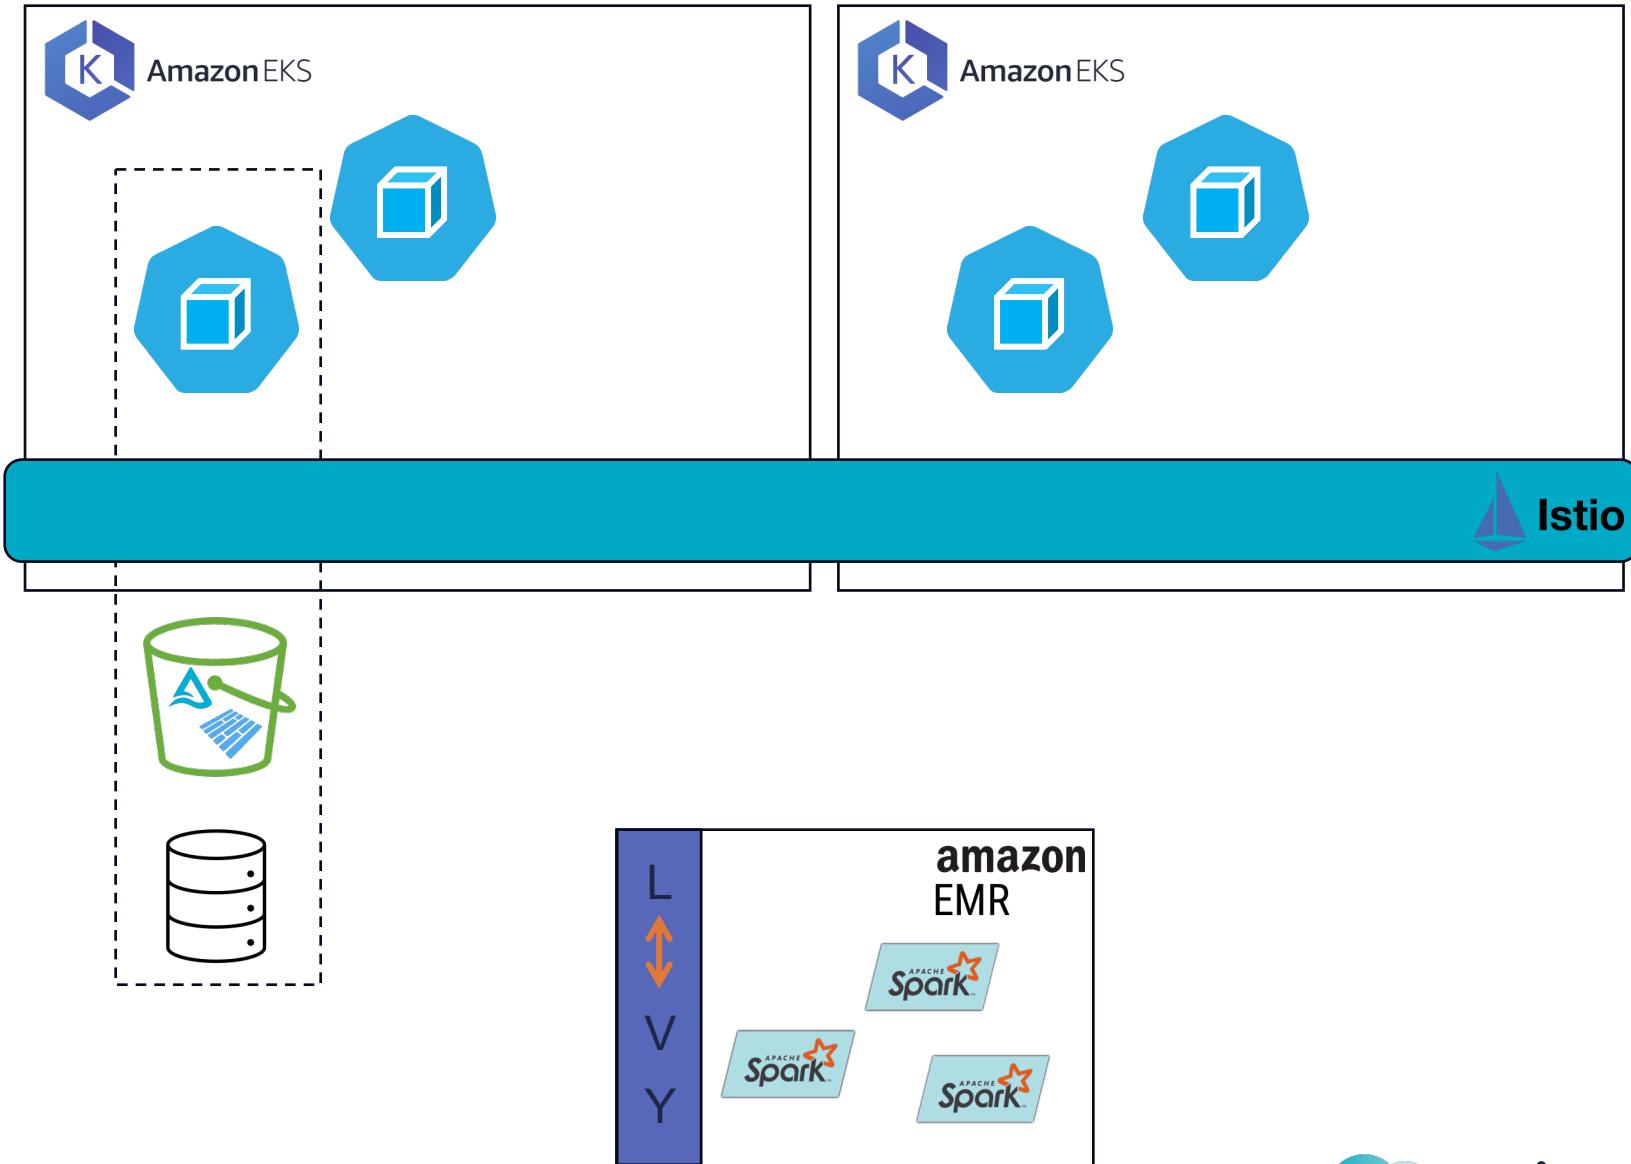
Technology Stack



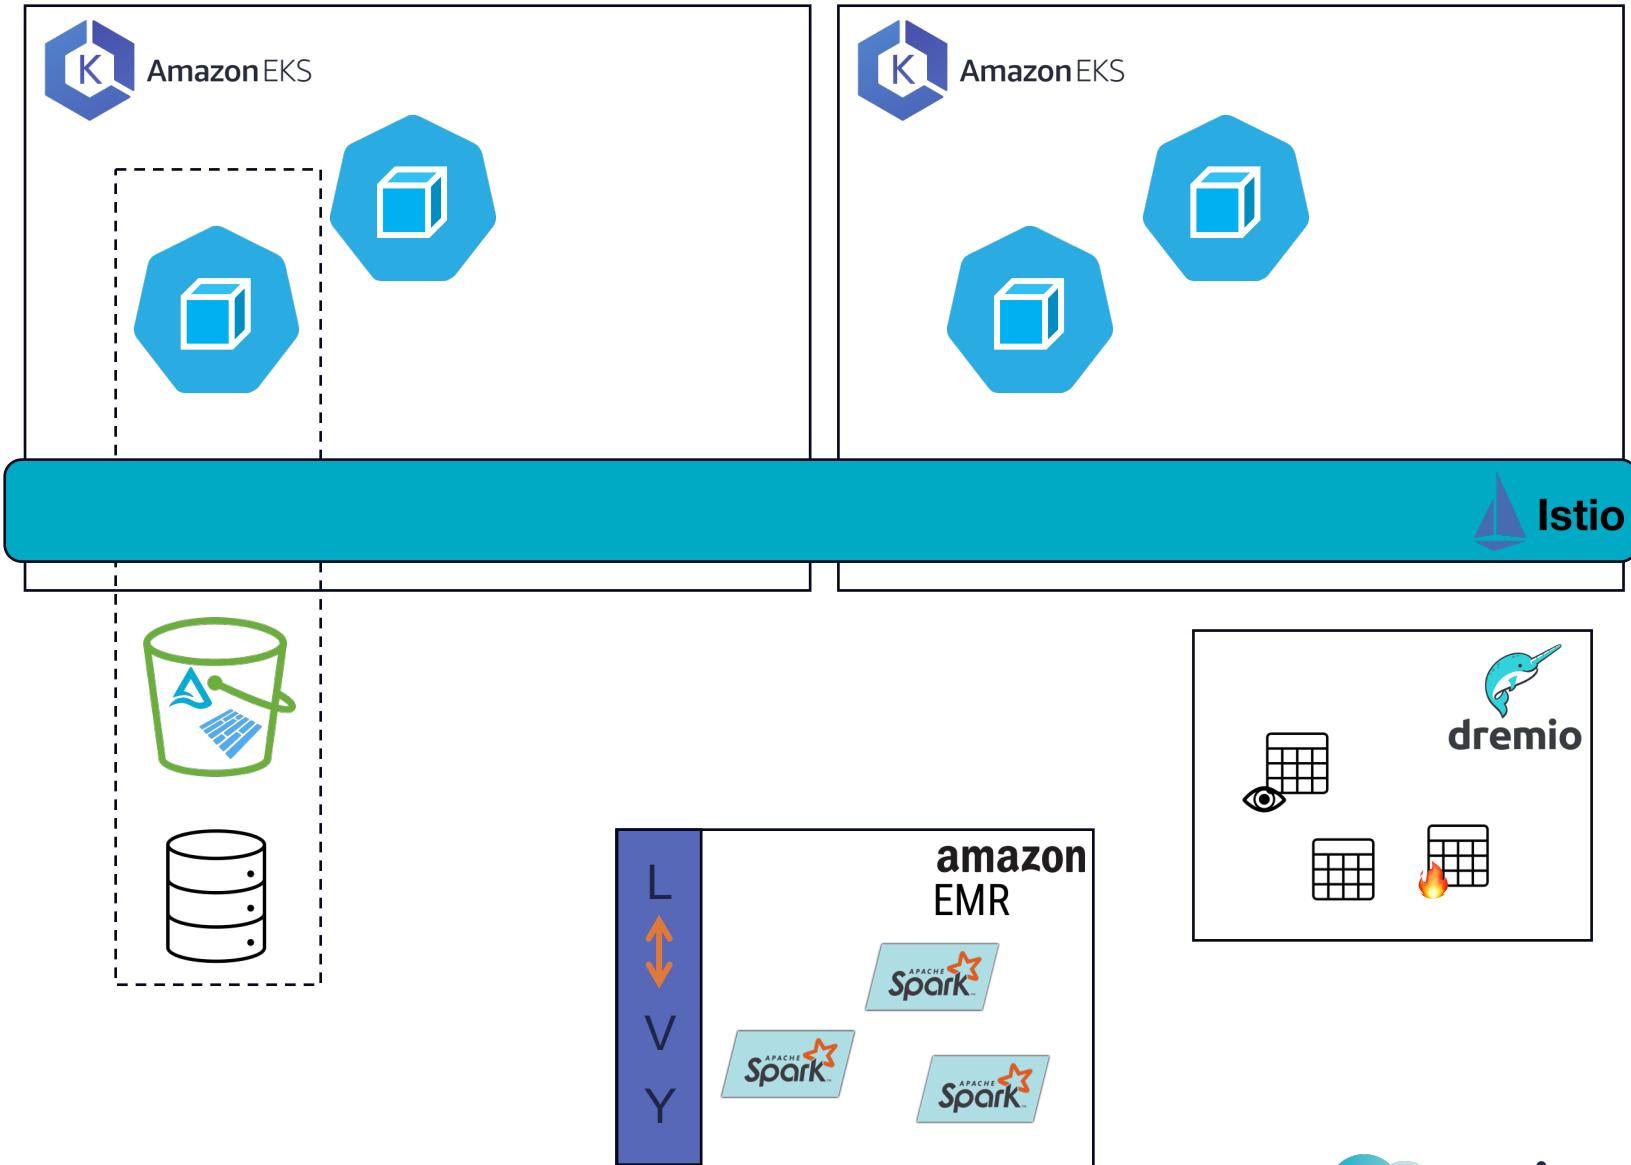
Technology Stack



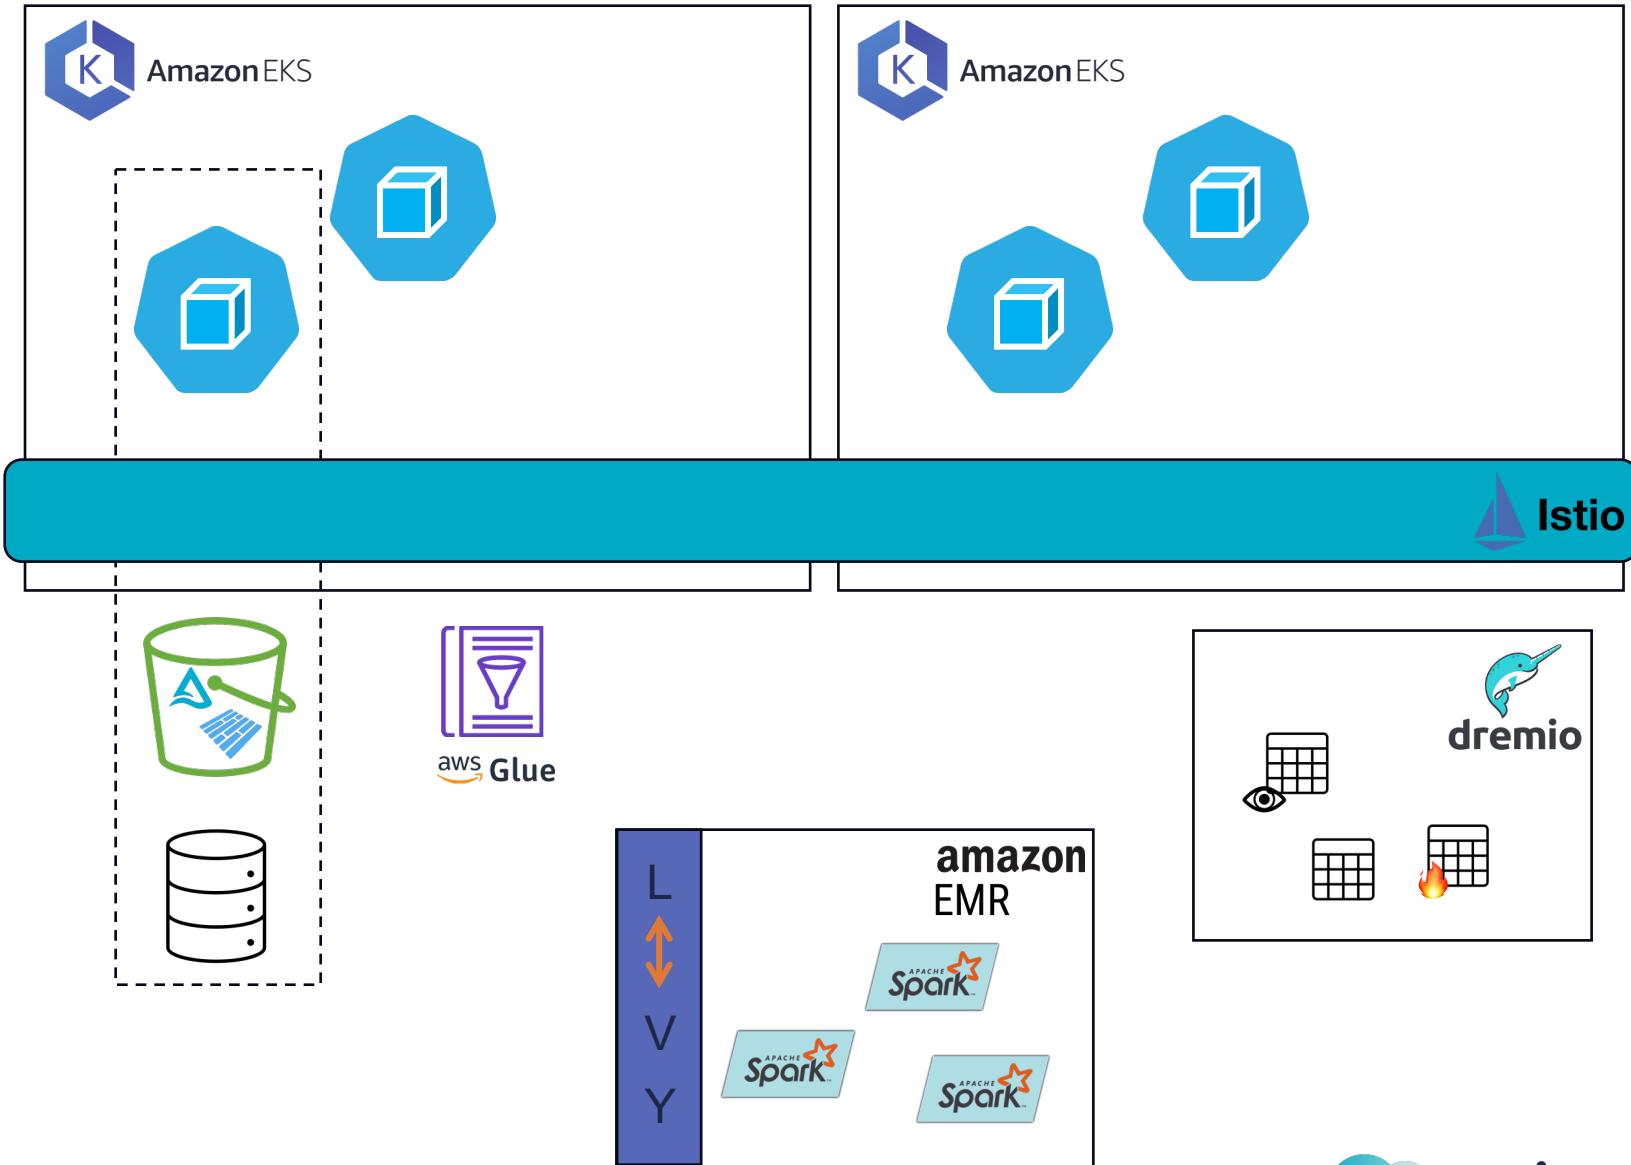
Technology Stack



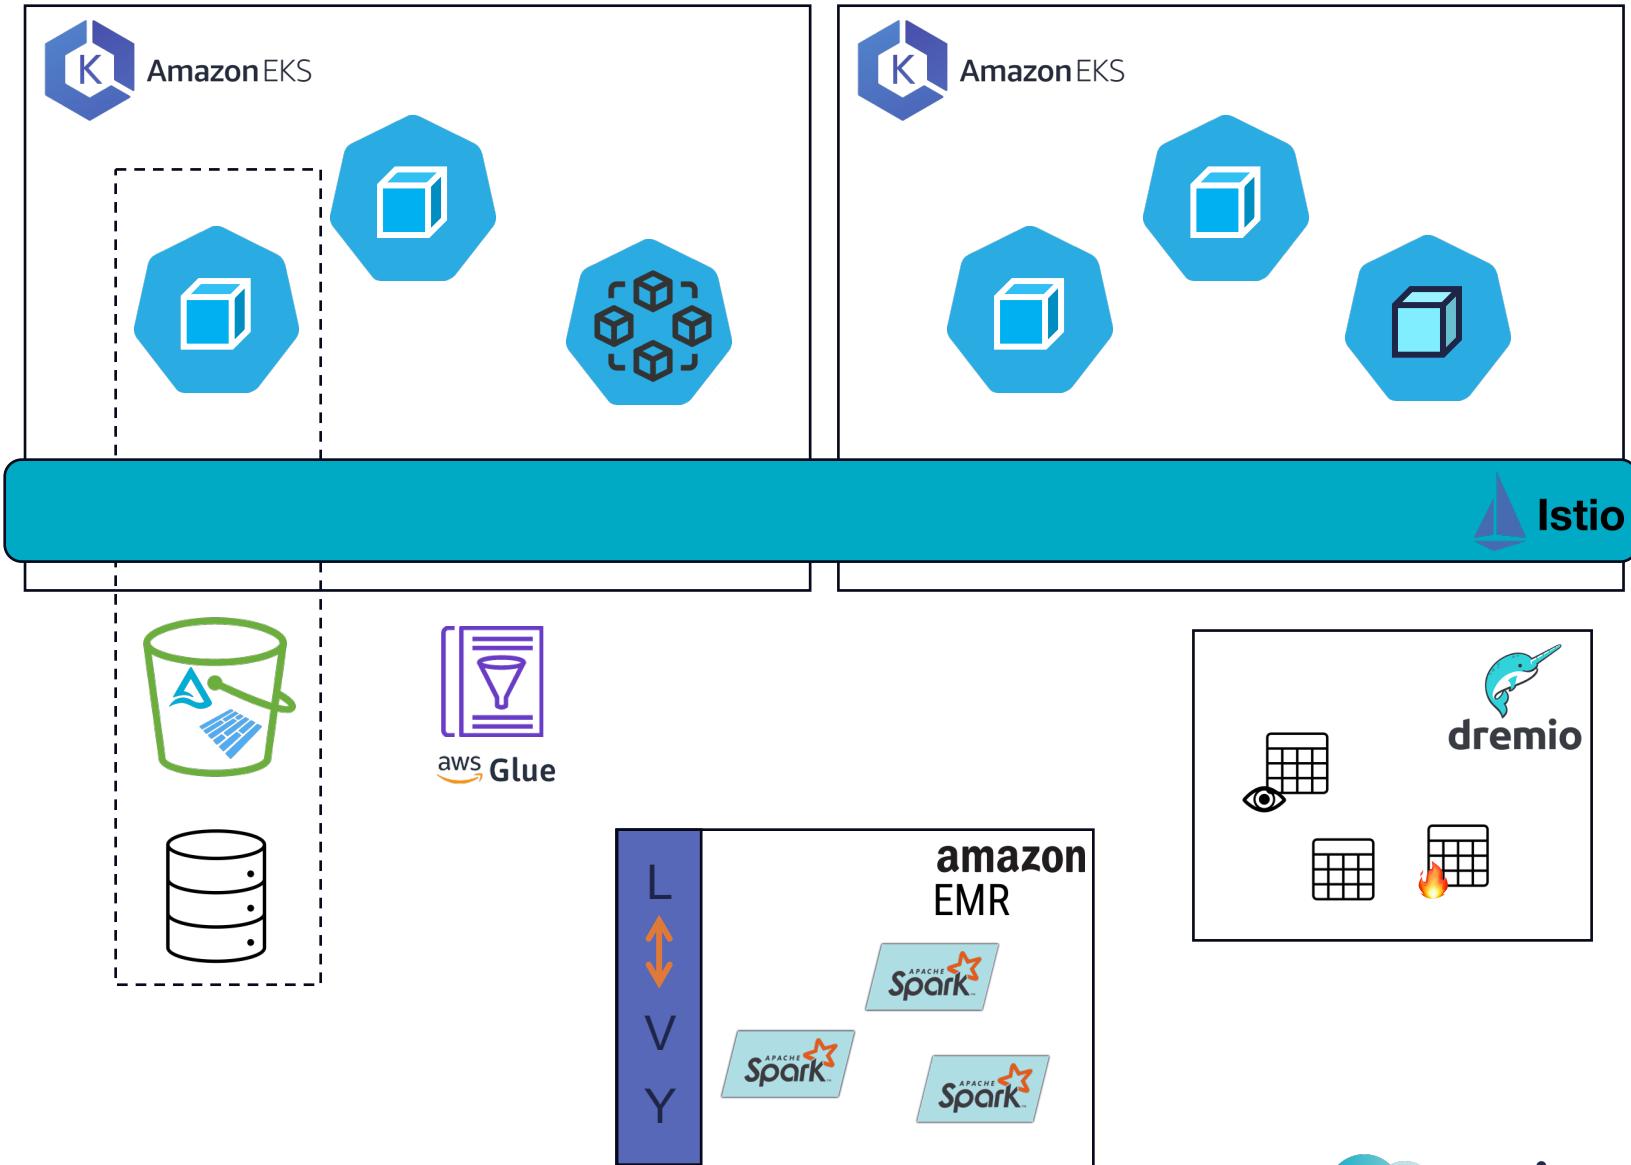
Technology Stack



Technology Stack



Technology Stack



Organization

Each domain team has an associated AWS IAM Role for each environment (dev/test/prod)

Each service has its own s3 bucket and operational DB

A central catalog is the entry-point for governance

Datasets are declared as Glue tables

Data for analytical purposes can be exposed in one of 3 ways:

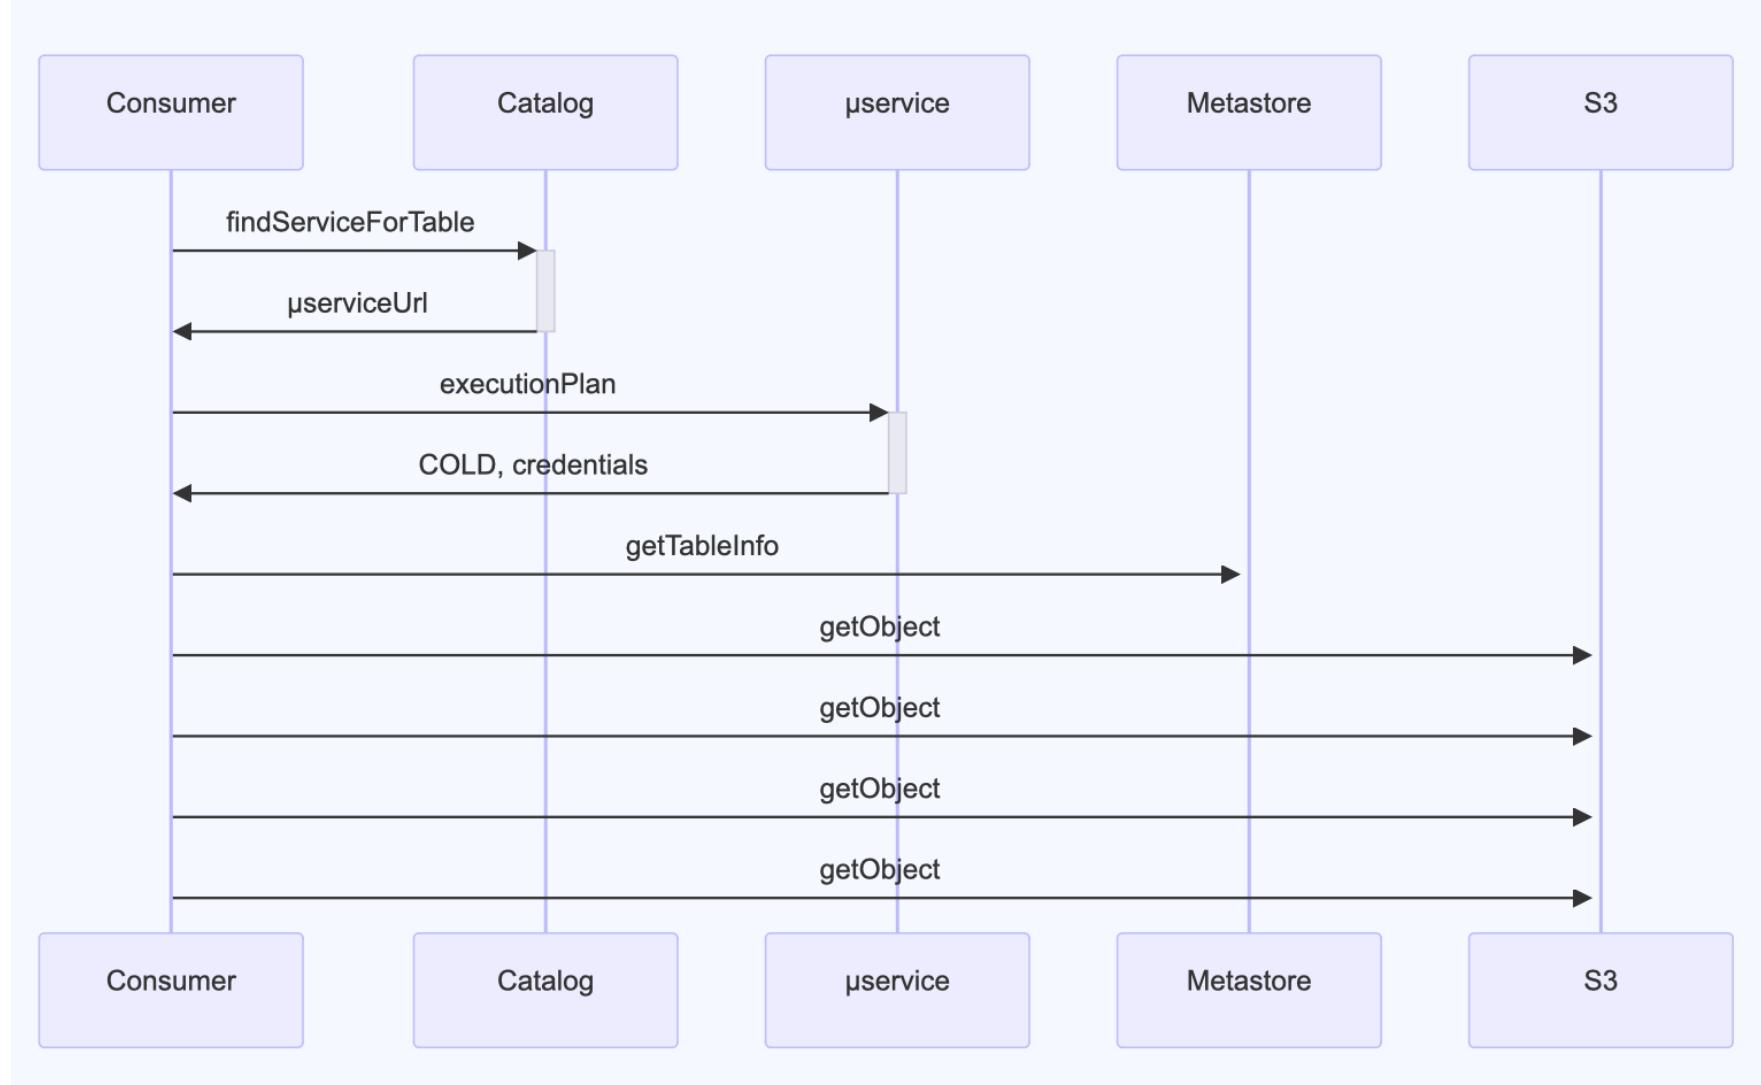
- ❄️ Parquet or Delta Lake tables on s3
- 🔥 json/http with a custom pre-defined protocol

Everything is provisioned through CloudFormation

Everything is API-first

Cold Data

(Parquet/Delta Lake on s3)



Hot 🔥 Data (*http/json*)

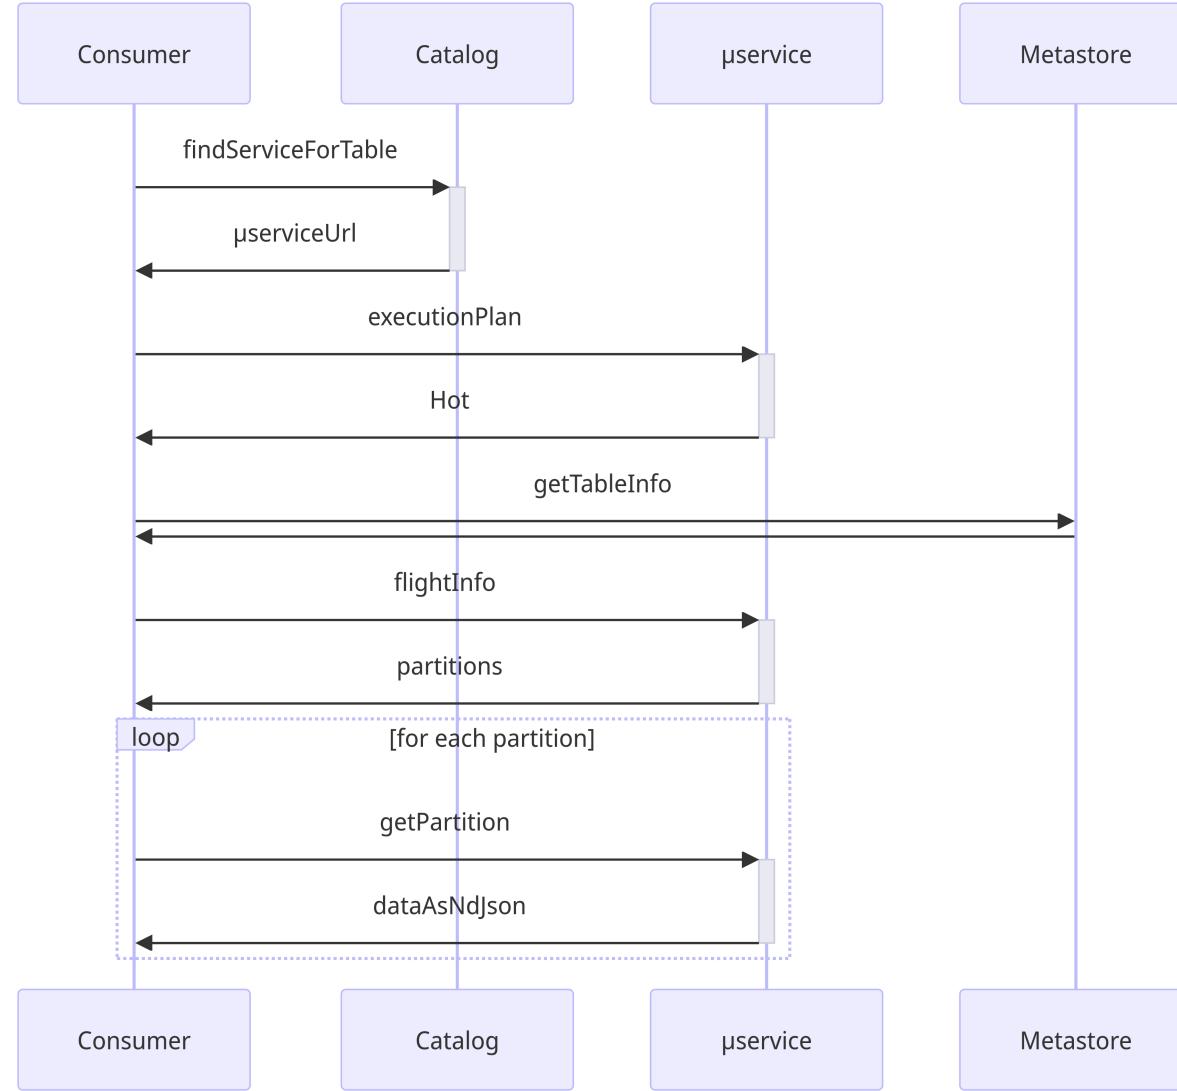
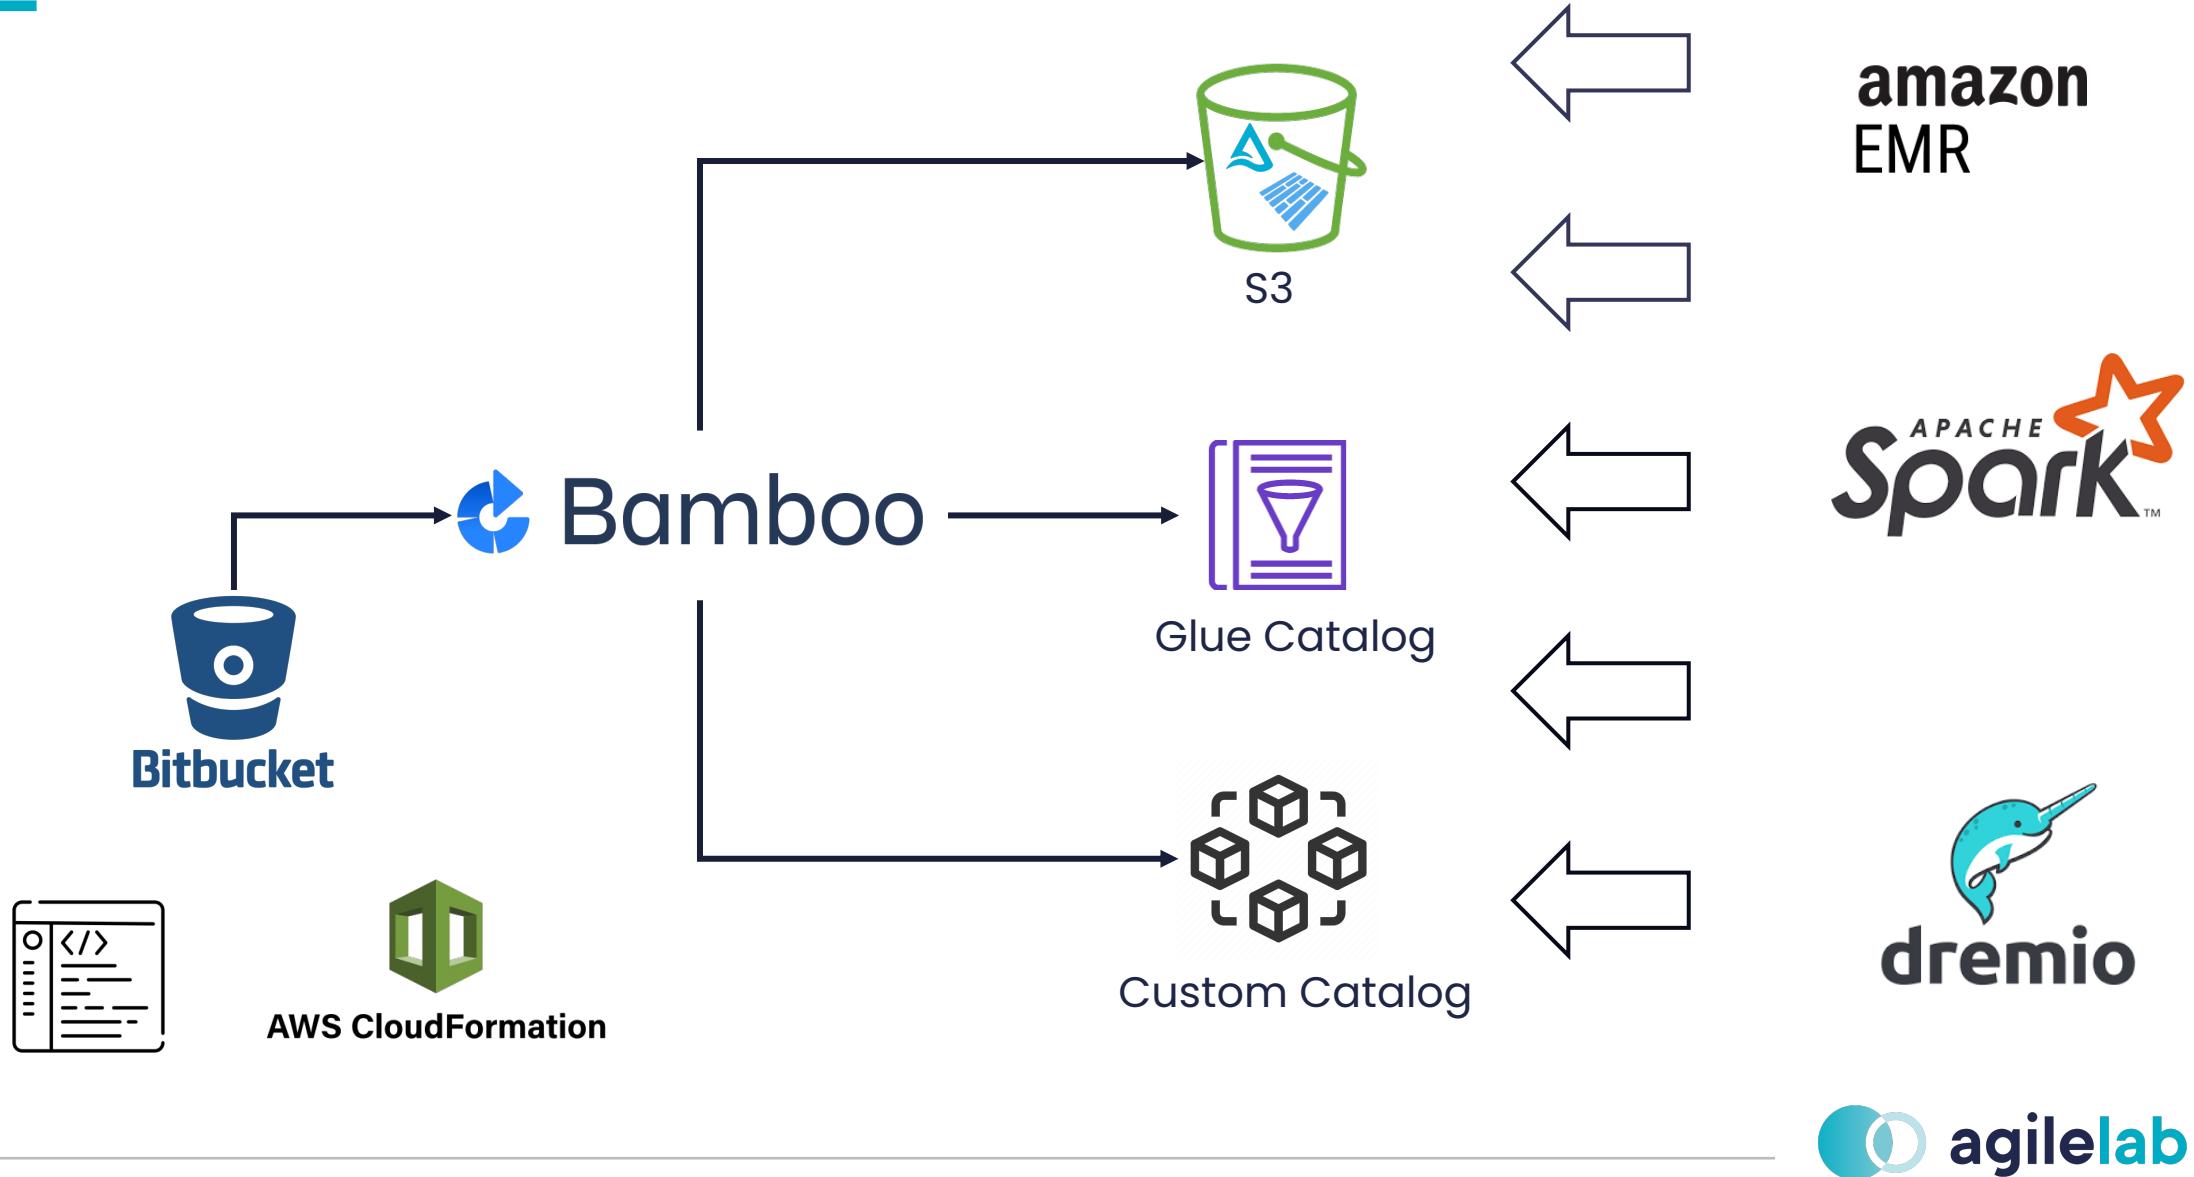
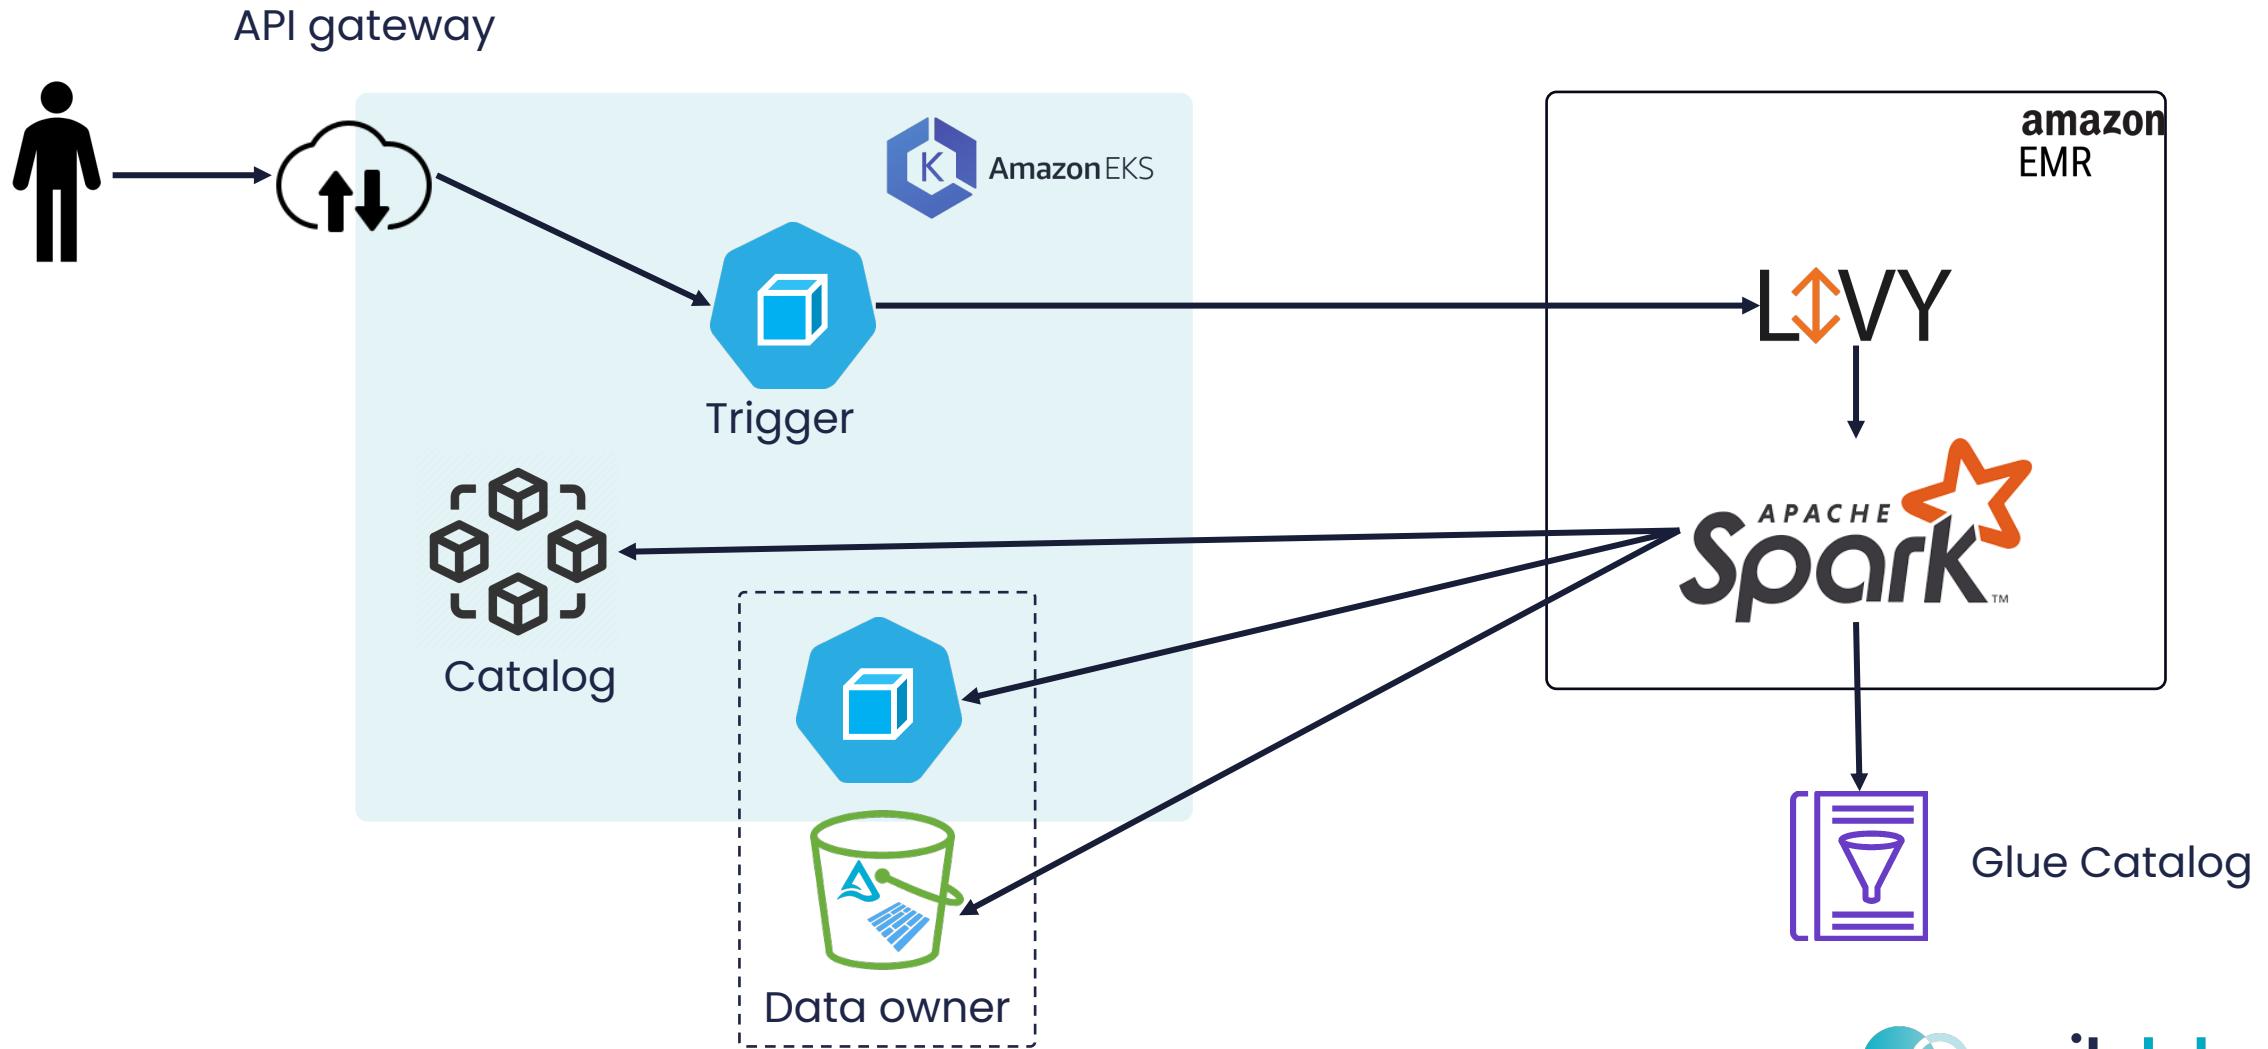


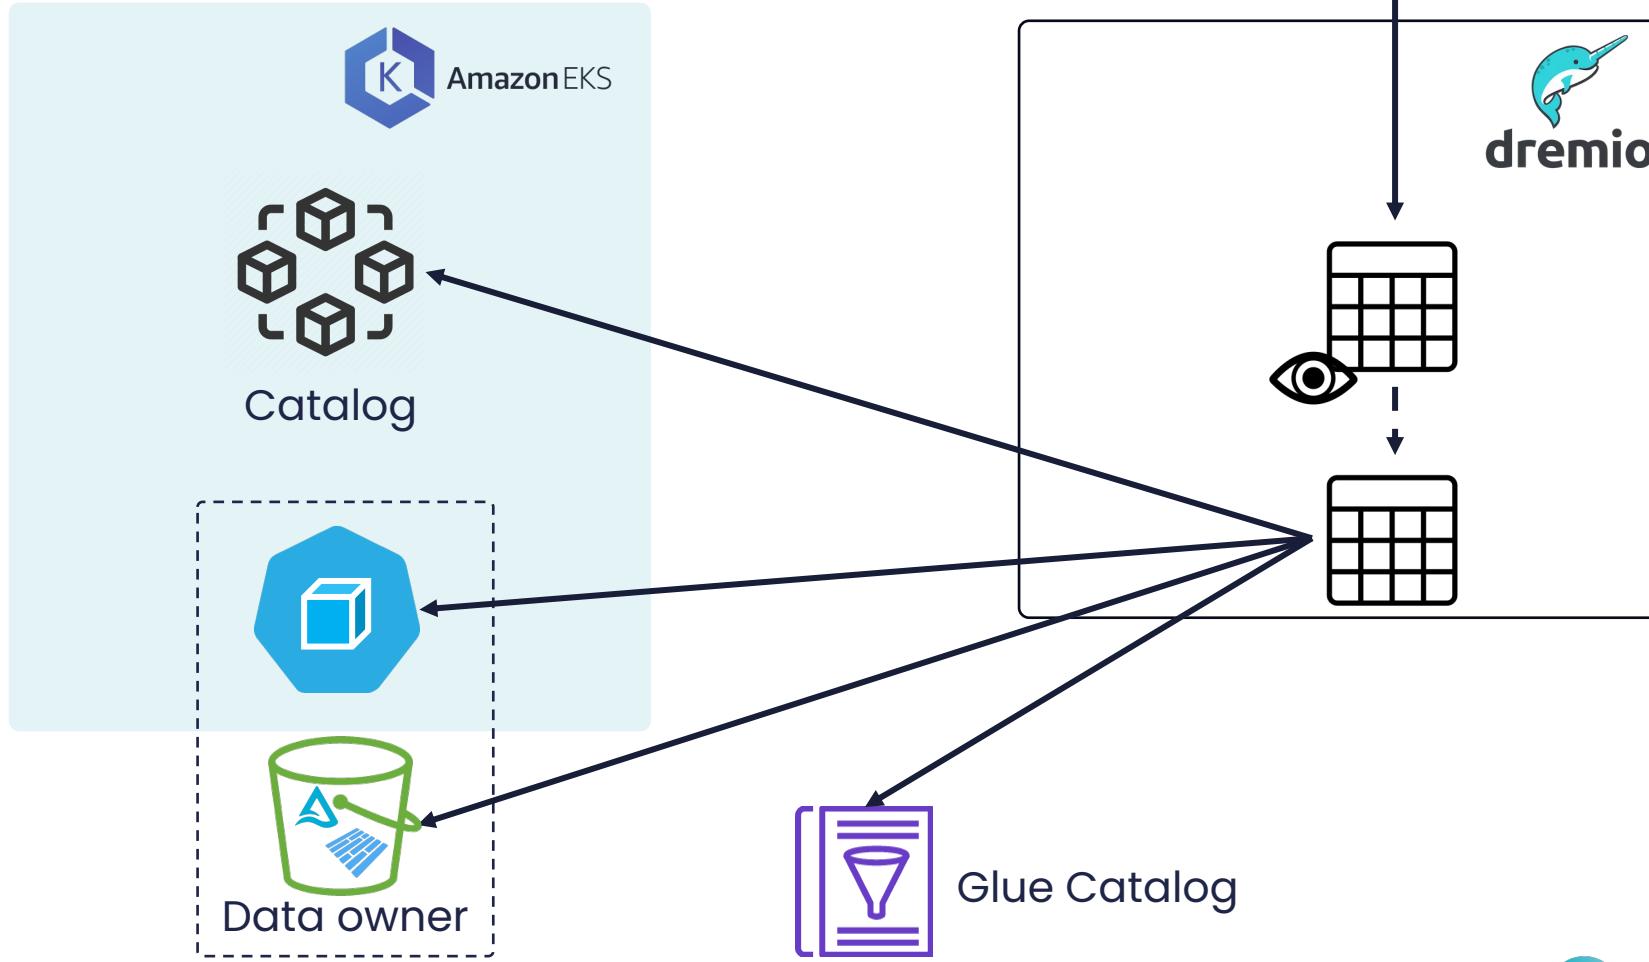
Table creation



Job submission



Dremio Access Model



**Let's dive into
Multi-cloud**

What's multicloud?



Definition

multiple cloud computing services in a single architecture

Key Points



Goals

Transparent allocation on  aws and  Azure of:

Compute



Storage



Already sorted out



«If I have seen further it is by standing on the shoulders of giants»

Isaac Newton

Http communication across different Kubernetes clusters on
different cloud providers (backbone on on-premise data center)

Challenges

No open source Spark Glue connector



EMR is not available out of aws

amazon
EMR

Provide a non-breaking upgrade



Metastore



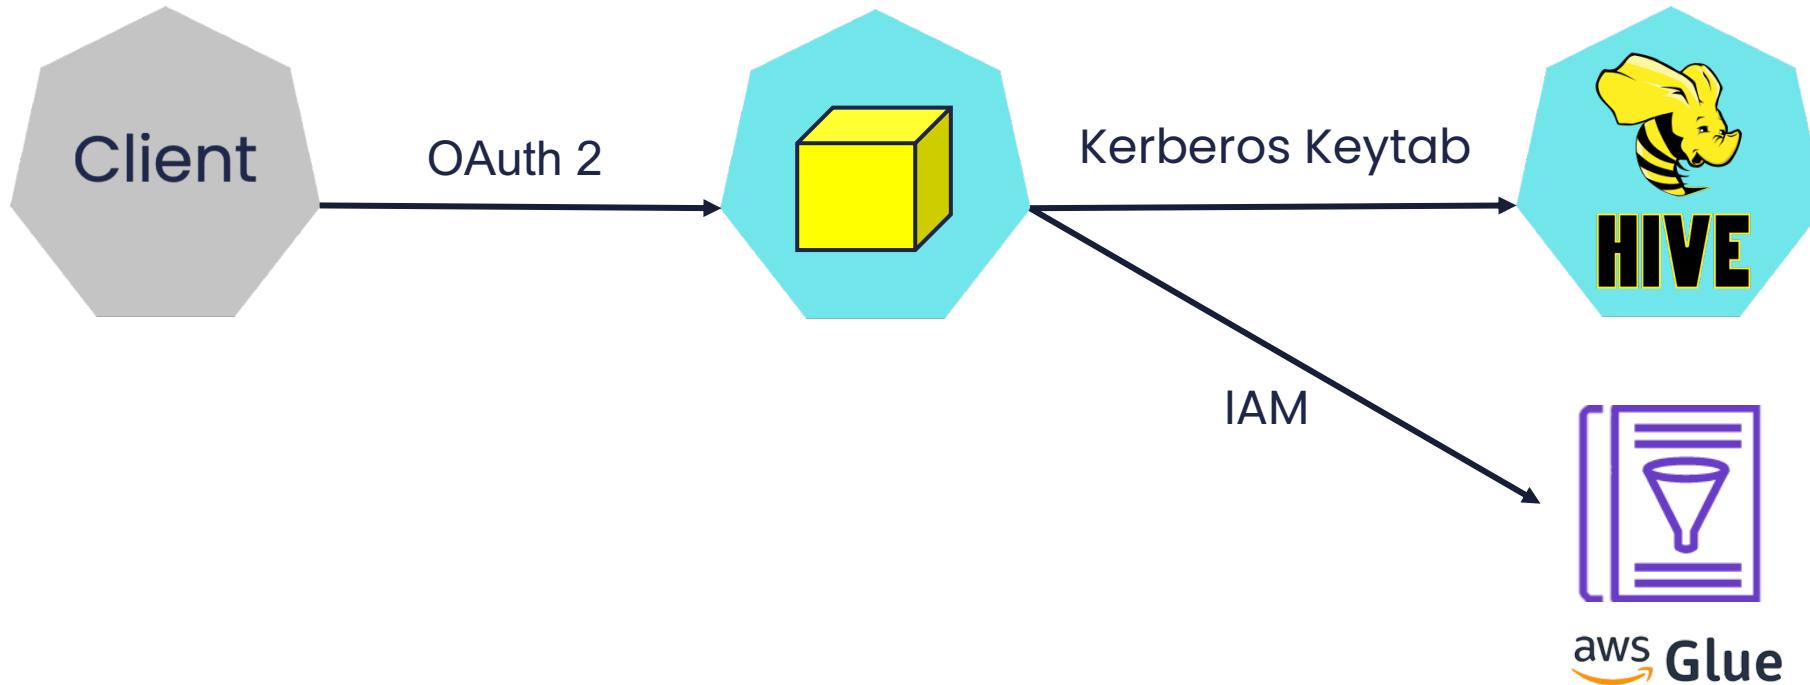
Challenges:

- 💀 Authentication is based on Kerberos
- ✨ Spark 2 can't easily switch between catalogs

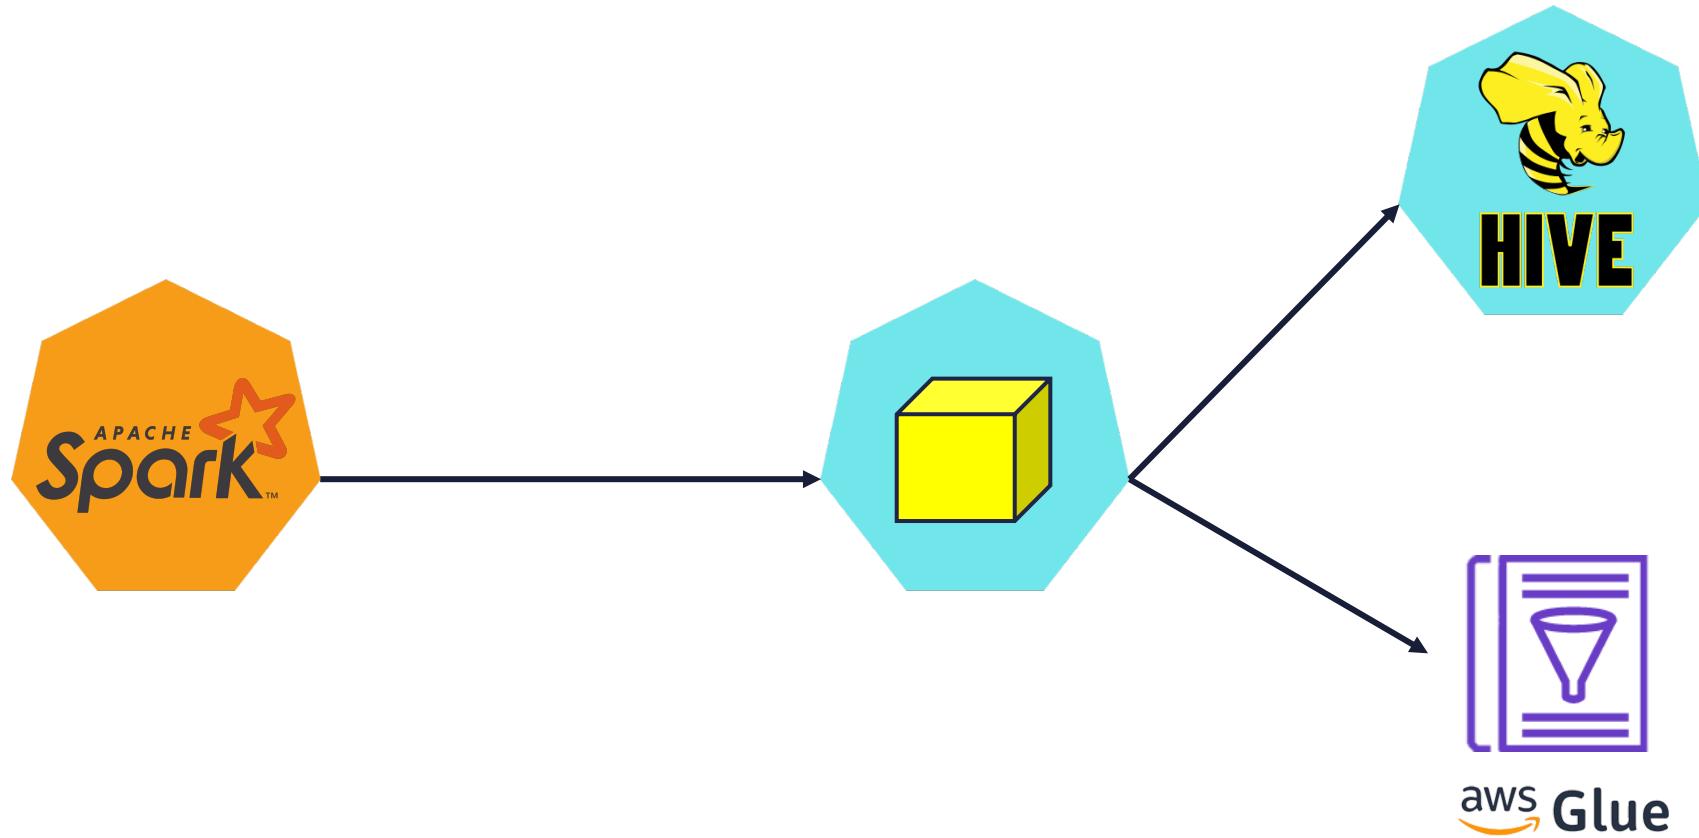
Kerberos

Hundreds of developers struggling with Kerberos?

Not on my watch



Spark single catalog limitation



Resource Manager for Spark



✓ Deployable in any cloud

✗ Huge pain to operationalise

Resource Manager for Spark

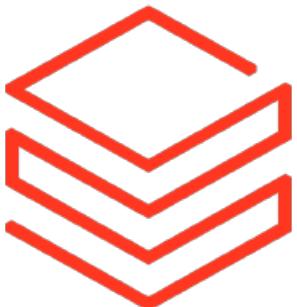


CLOUDERA

- ✓ Deployable in any cloud
- ✓ Easier to operationalise

- ✗ Quite hard to integrate across instances

Resource Manager for Spark



- ✓ Deployable in any cloud
- ✓ Easier to operationalise
- ✓ Federation ease integration

✗ Vendor lock-in

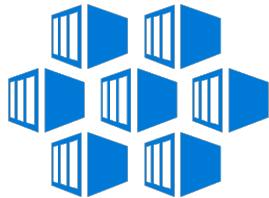
Resource Manager for Spark



- ✓ Deployable in any cloud
- ✓ No vendor lock-in
- ✓ Good isolation
- ✓ Huge configurability

- ✗ Huge pain to operationalise
- ✗ Maintain your Spark “distro”

Resource Manager for Spark



- ✓ Deployable in any cloud
- ✓ Easier to operationalise
- ✓ No vendor lock-in
- ✓ Good isolation
- ✓ Huge configurability

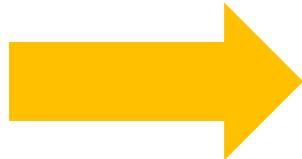
✗ Maintain your Spark “distro”

Smooth migration – evil master plan

1. Hide low level details
2. Introduce concept of deployment options
3. Build multi-cloud support
4. Add Azure deployment options
5. Gently roll out upgrades
6. Profit 🤞

Step 1.1 – Hide storage and table provisioning

Cloudformation



Cloud agnostic descriptor

Step 1.1 – Hide storage and table provisioning

```
version: v3
type: parallelread
entityName: customer
entityVersion: v1
description: customer data from the financial system
coldStorage:
  format: parquet
  location:
    bucket: finance_customer_(){}
    prefix: bronze
    provider: aws|azure
  hasHistory: false
  partitioningColumns: [country]
columns:
  - name: fullName
    type: string
      - dgItaly1
  - name: address
    type: string
    optional: true
```

Step 1.1 – Hide storage and table provisioning

```
version: v3
type: parallelread
entityName: customer
entityVersion: v1
description: customer data from the financial system
coldStorage:
  format: parquet
  location:
    bucket: finance_customer_(){}
    prefix: bronze
    provider: aws | azure
  hasHistory: false
  partitioningColumns: [country]
columns:
  - name: fullName
    type: string
      - dgItaly1
  - name: address
    type: string
    optional: true
```

Step 1.2 – hide Spark low level details



Step 1.2 – hide Spark low level details



pre-built Templates

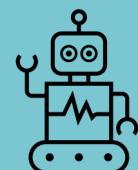
Step 2.1 – Deployment options



Introduce batch job deployment descriptor with deployment target



The project skeleton/template is auto-configured from the deployment descriptor



The CI/CD pipeline takes care of setting everything up

Step 2.1 – Deployment options

```
1 specVersion: "v1"
2 name: "helloworld"
3 sources:
4   ...
5     domain: "Finance"
6         name: "customer"
7     ...
8     domain: "Ecommerce"
9         name: "orders"
10        version: "v2"
11 sink:
12   ...
13     domain: "Marketing"
14     name: "most_spending_customers"
15     version: "v1"
16 framework: "#technology-specific-configuration"
17     kind: spark-2.4.4-emr6-livy-scala
18     cluster: emr1
19     entrypointClass: "io.contoso.app.HelloWorld"
20     dependencies: []
21     configurations:
22       ...
23       spark.task.maxFailures: 6
24       ...
25       spark.dynamicAllocation.enabled: false
```

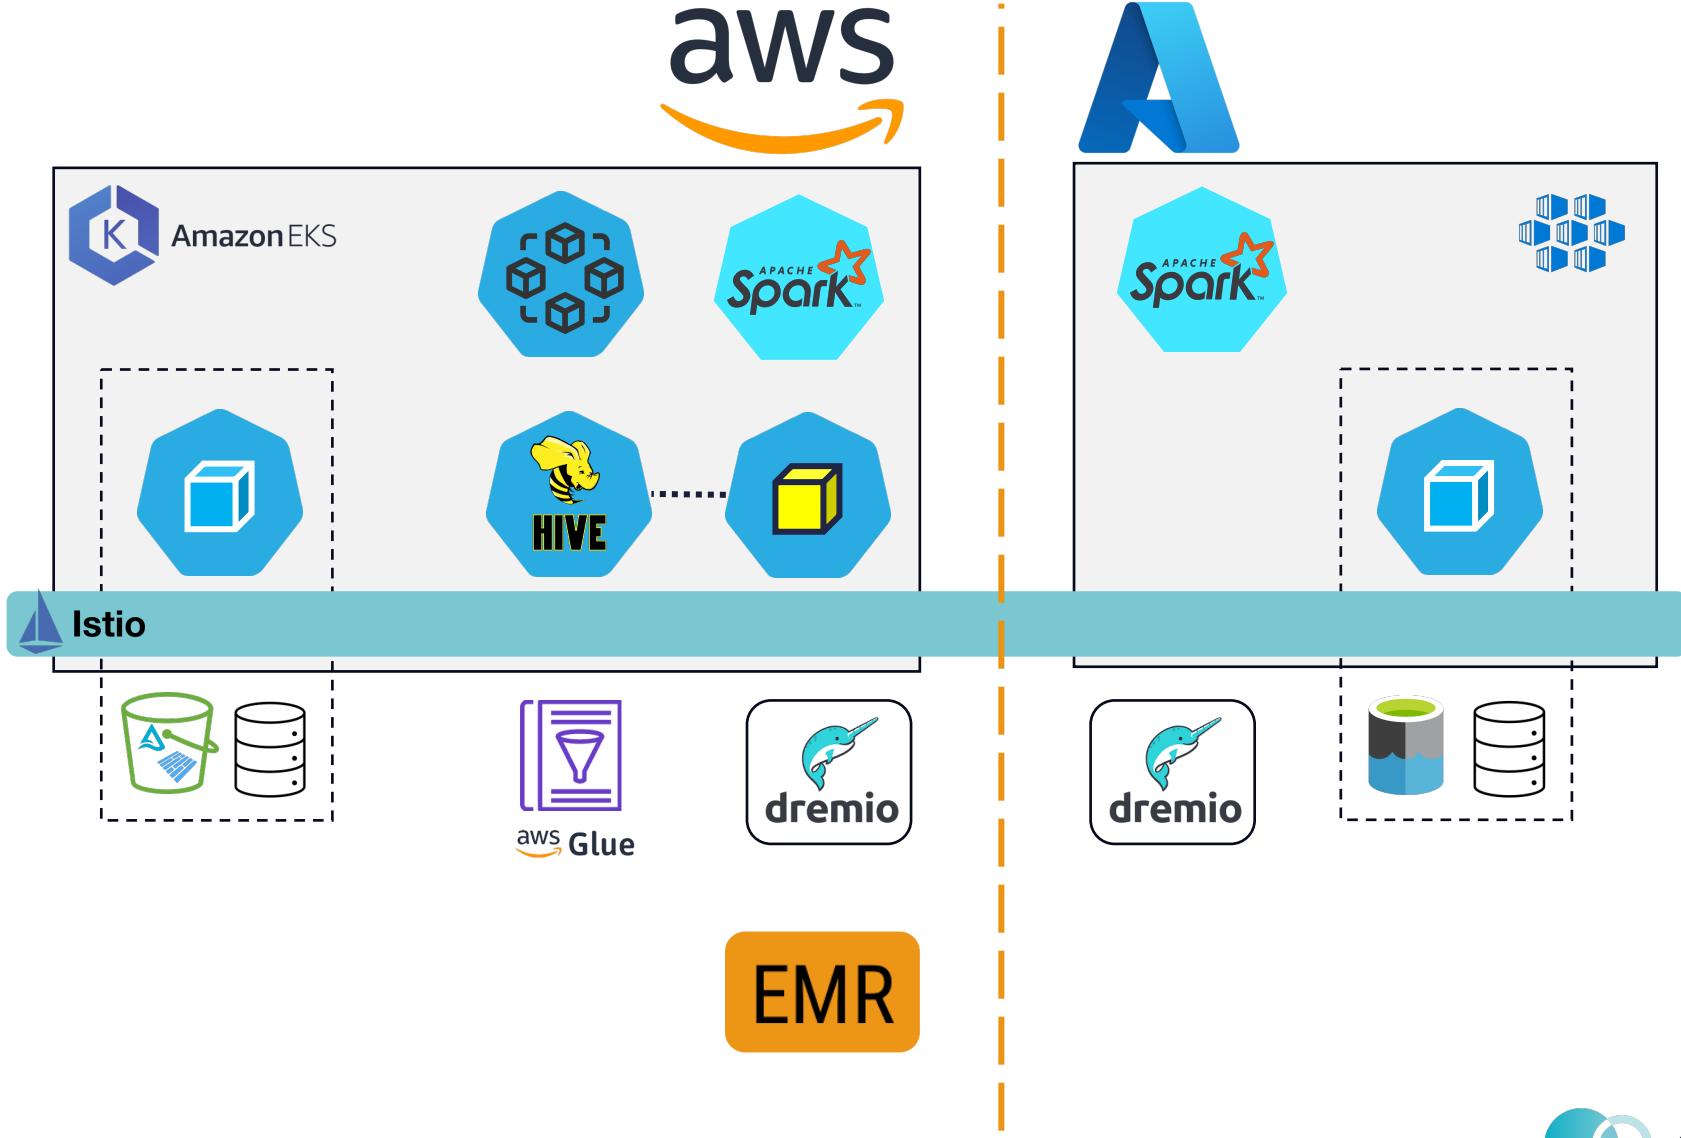
Step 3 – Build azure support

- Extend our sdk to support Azure/ADLS
- Prepare base container images
- Hide s3/adls setup on SDK side
- Setup spark context with custom catalog

Step 4 – Add azure support

```
specVersion: v1
name: "helloworld"
sources:
  - domain: Finance
    name: customer
    version: v1
  - domain: Ecommerce
    name: orders
    version: v2
sink:
  domain: Marketing
  name: most_spending_customer
  version: v1
framework: # technologyspecific configuration
  kind: spark-3.4.1-k8s-livy-scala
  cluster: azure1
entrypointClass: io.contoso.app.HelloWorld
dependencies: [ ]
```

Full picture



Gently roll-out

- We don't FORCE anyone to update to Spark newest version
- New applications run on Kubernetes by default
- Applications running on EMR can't access datasets stored on Azure
- So we give them a good reason to update!

Roll out status – 6 months in

Thanks to the platform custom catalog we know that:



multi-cloud



New app on Azure
(3% overall)

Lessons learnt



Abstracting infra is good way of handling a transition

Truly open standards are there and help a lot

You can/should create your internal standards (when needed)

You **need** a Platform Team

Data has gravity (egress traffic is expensive)

Next steps – caching

There are only two hard things in Computer Science: cache invalidation and naming things.

Phil Karlton

Next steps – smart scheduler

My favourite subject at school was avoiding unnecessary work

Prince Philip

Next steps – true independence



Independence is not freedom, it's a responsibility

Thank you!



x

x

Questions?



www.agilelab.it



@agilelab_official



linkedin.com/company/agile-lab



@agilelabsrl